

Open Domain Event Extraction from Twitter

Abstract

Tweets are the most up-to-date and inclusive stream of information and commentary on current events, but they are also fragmented and noisy, motivating the need for systems that can extract, aggregate and categorize important events. Previous work on extracting structured representations of events has focused largely on newswire text; Twitter’s unique characteristics present new challenges and opportunities for open-domain event extraction. This paper describes TWICAL—the first open-domain event-extraction and categorization system for Twitter. We demonstrate that accurate open-domain event extraction from Twitter is indeed feasible. In addition we present an approach to categorizing the extracted events based on latent variable models. By leveraging large volumes of unlabeled data, our latent-variable model achieves a 31% F1 increase over a supervised baseline. A continuously updating demonstration of our system can be viewed at http://anonymized_for_review.

1 Introduction

Social networking sites such as Facebook and Twitter present the most up-to-date information and buzz about current events. Yet the number of tweets posted daily has recently exceeded two-hundred million, many of which are either redundant (Zanzotto et al., 2011), or of limited interest, leading to information overload.¹ Clearly, we can benefit from more

structured representations of events that are *synthesized* from individual tweets.

Previous work in event extraction (Grishman and Sundheim, 1996; Allan et al., 1998; Yang et al., 1998; Pustejovsky et al., 2003; Doddington et al., 2004; Chambers and Jurafsky, 2011) has focused mostly on news articles, as historically this genre of text has been the best source of information on current events. In the meantime, social networking sites such as Facebook and Twitter have become an important complementary source of such information. While status messages contain a wealth of useful information, they are very disorganized motivating the need for automatic extraction, aggregation and categorization.

Extracting useful structured representations of events from this disorganized corpus of noisy text is a challenging problem. On the other hand, individual tweets are short and self-contained and therefore don’t require complex discourse processing as is the case for texts containing narratives. In this paper we demonstrate that **open-domain event extraction** from Twitter is indeed feasible, for example our highest-confidence extracted future events are 90% accurate as demonstrated in §6.

Twitter has several characteristics which present unique *challenges* and *opportunities* for the task of open-domain event extraction.

Challenges: Twitter users frequently mention mundane events in their daily lives (such as what they ate for lunch) which are only of interest to their immediate social network. In contrast, if an event is mentioned in newswire text, it is safe to assume it is of general importance. Indi-

¹<http://blog.twitter.com/2011/06/200-million-tweets-per-day.html>

vidual tweets are also very terse, often lacking sufficient context to categorize them into topics of interest (e.g. SPORTS, POLITICS, PRODUCTRELEASE etc...). Further because Twitter users can talk about whatever they choose, it is unclear in advance which set of event types are appropriate. Finally, tweets are written in an informal style causing NLP tools designed for edited texts to perform extremely poorly.

Opportunities: The short and self-contained nature of tweets means they have very simple discourse and pragmatic structure, issues which still challenge state-of-the-art NLP systems. For example in newswire, complex reasoning about relations between events (e.g. *before* and *after*) is often required to accurately relate events to temporal expressions (Mani et al., 2006; Chambers et al., 2007). The volume of Tweets is also much larger than the volume of news articles, so redundancy of information can be exploited more easily.

To address Twitter’s noisy style, we follow recent work on NLP in noisy text (Ritter et al., 2011; Liu et al., 2011; Gimpel et al., 2011), annotating a corpus of Tweets with events,² which is then used as training data for sequence-labeling models to identify event mentions in millions of messages.

Because of the terse, sometimes mundane, but highly redundant nature of tweets, we were motivated to focus on extracting an aggregate representation of events which provides additional context for tasks such as event categorization, and also filters out mundane events by exploiting redundancy of information. We propose identifying important events as those whose mentions are strongly associated with references to a unique date as opposed to dates which are evenly distributed across the calendar.

Twitter users discuss a wide variety of topics, making it unclear in advance what set of event types are appropriate for categorization. To address the diversity of events discussed on Twitter, we introduce a novel approach to discovering important event types and categorizing aggregate events within a new domain.

²See the supplementary materials.

Entity	Event Phrase	Date	Type
Steve Jobs	died	10/6/11	DEATH
iPhone	announcement	10/4/11	PRODUCTLAUNCH
GOP	debate	9/7/11	POLITICALEVENT
Amanda Knox	verdict	10/3/11	TRIAL

Table 1: Examples of events extracted by TWICAL.

Supervised or semi-supervised approaches to event categorization would require first designing annotation guidelines (including selecting an appropriate set of types to annotate), then annotating a large corpus of events found in Twitter. This approach has several drawbacks, as it is apriori unclear what set of types should be annotated; a large amount of effort would be required to manually annotate a corpus of events while simultaneously refining annotation standards.

We propose an approach to open-domain event categorization based on latent variable models that uncovers an appropriate set of types which match the data. The automatically discovered types are subsequently inspected to filter out any which are incoherent and the rest are annotated with informative labels;³ examples of types discovered using our approach are listed in figure 2. The resulting set of types are then applied to categorize hundreds of millions of extracted events without the use of any manually annotated examples. By leveraging large quantities of unlabeled data, our approach results in a 31% improvement in F₁ score over a supervised baseline which uses the same set of types.

2 System Overview and Background

TWICAL extracts a 4-tuple representation of events which includes a named entity, event phrase, calendar date, and event type (see Table 1). This representation was chosen to closely match the way events are typically mentioned in Twitter.

An overview of the various components of our system for extracting events from Twitter is presented in Figure 1. Given a raw stream of tweets, our system extracts named entities in association with event phrases and unambiguous dates which are involved in significant events. First the tweets are POS

³This annotation and filtering takes minimal effort. One of the authors spent roughly 30 minutes inspecting and annotating the automatically discovered event types.

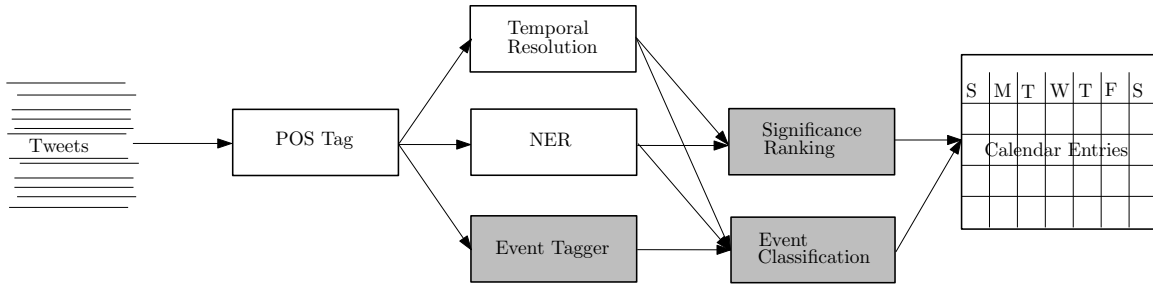


Figure 1: Processing pipeline for extracting events from Twitter. New components developed as part of this work are shaded in grey.

tagged, then named entities and event phrases are extracted, temporal expressions resolved, and the extracted events are categorized into types. Finally we measure the strength of association between each named entity and date based on the number of tweets they co-occur in, in order to determine whether an event is significant.

NLP tools, such as named entity segmenters and part of speech taggers which were designed to process edited texts (e.g. news articles) perform very poorly when applied to Twitter text due to its noisy and unique style. To address these issues, we utilize a named entity tagger and part of speech tagger trained on in-domain Twitter data presented by Ritter et. al. (2011).⁴ We also develop an event tagger trained on in-domain annotated data as described in §3.

In addition to extracting events and related named entities, we also need to extract when they occur. In general there are many different ways users can refer to the same calendar date, for example “next Friday”, “August 12th”, “tomorrow” or “yesterday” could all refer to the same day, depending on when the tweet was written. To resolve temporal expressions we make use of TempEx (Mani and Wilson, 2000), which takes as input a reference date, some text, and parts of speech (from the Ritter et. al. Twitter-trained POS tagger) and marks temporal expressions with unambiguous calendar references. Although this mostly rule-based system was designed for use on newswire text, we find its precision on Tweets (94% - estimated over as sample of 268 extractions) is sufficiently high to be useful for our purposes. TempEx’s high precision on

Tweets can be explained by the fact that most temporal expressions are relatively unambiguous. Although there appears to be room for improving the recall of temporal extraction on Twitter by handling noisy temporal expressions (for example see Ritter et. al. (2011) for a list of over 50 spelling variations on the word “tomorrow”), we leave adapting temporal extraction to Twitter as potential future work.

3 Extracting Event Mentions

In order to extract event mentions from Twitter’s noisy text, we first annotate a corpus of tweets, which is then used to train sequence models to extract events. While we apply an established approach to sequence-labeling tasks in noisy text (Ritter et al., 2011; Liu et al., 2011; Gimpel et al., 2011), this is the first work to extract event-referring phrases in Twitter.

Event phrases can consist of many different parts of speech as illustrated in the following examples:

- **Verbs:** Apple to *Announce* iPhone 5 on October 4th?! YES!
- **Nouns:** iPhone 5 *announcement* coming Oct 4th
- **Adjectives:** WOOOHOO *NEW* IPHONE TODAY! CAN’T WAIT!

These phrases provide important context, for example extracting the entity, **Steve Jobs** and the event phrase **died** in connection with October 5th, is much more informative than simply extracting **Steve Jobs**. In addition, event mentions are helpful in upstream tasks, for example categorizing events into types, as described in §4.

⁴Available at http://github.com/aritter/twitter_nlp.

	precision	recall	F1
TWICAL-EVENT	0.56	0.74	0.64
No POS	0.48	0.70	0.57
Timebank	0.24	0.11	0.15

Table 2: Precision and recall at event phrase extraction. All results are reported using 4-fold cross validation over the 1,000 manually annotated tweets (about 19K tokens). We compare against a system which doesn’t make use of features generated based on the Ritter et. al. POS Tagger, in addition to a system trained on the Timebank corpus which uses the same set of features.

In order to build a tagger for recognizing events, we annotated 1,000 tweets (19,484 tokens) with event phrases, following annotation guidelines similar to those developed for annotating the EVENT tags in Timebank (Pustejovsky et al., 2003).⁵ We treat the problem of recognizing event triggers as a sequence labeling task, using Conditional Random Fields for learning and inference (Lafferty et al., 2001), and making use of contextual, dictionary, and orthographic features. We also include features based on the Twitter-tuned POS tagger of Ritter et. al. (2011), and dictionaries of event terms gathered from WordNet by Sauri et al. (2005).

The precision and recall at segmenting event phrases are reported in Table 2. Our classifier, TWICAL-EVENT, obtains an F-score of 0.64. To demonstrate the need for in-domain training data, we compare against a baseline of training our system on the Timebank corpus.

4 Classification of Event Types

To categorize the extracted events into types we propose an approach based on latent variable models which infers an appropriate set of event types to match our data, and also classifies events into types by leveraging large amounts of unlabeled data.

Supervised or semi-supervised classification of event categories is problematic for a number of reasons. First, it is *a priori* unclear which categories are appropriate for Twitter. Secondly, a large amount of manual effort is required to annotate tweets with event types. Third, the set of important categories (and entities) is likely to shift over time, or within a

focused user demographic. Finally many important categories are relatively infrequent, so even a large annotated dataset may contain just a few examples of these categories, making classification difficult.

For these reasons we were motivated to investigate unsupervised approaches that will automatically induce event types which match the data. We adopt an approach based on latent variable models inspired by recent work on modeling selectional preferences (Ritter et al., 2010; Ó Séaghdha, 2010; Kozareva and Hovy, 2010; Van Durme and Gildea, 2009), and unsupervised information extraction (Yao et al., 2011; Chambers and Jurafsky, 2011).

Each event indicator phrase in our data, e , is modeled as a mixture of types. For example the event phrase “cheered” might appear as part of either a POLITICALEVENT, or a SPORTSEVENT. Each type corresponds to a distribution over named entities n involved in specific instances of the type, in addition to a distribution over dates d on which events of the type occur. Including calendar dates in our model has the effect of encouraging (though not requiring) events which occur on the same date to be assigned the same type. This is helpful in guiding inference, because distinct references to the same event should also have the same type.

The generative story for our data is based on Lin-kLDA (Erosheva et al., 2004), and is presented as Algorithm 1. This approach has the advantage that information about an event phrase’s type distribution is shared across it’s mentions, while ambiguity is also naturally preserved. In addition, because the approach is based on generative a probabilistic model, it is straightforward to perform many different probabilistic queries about the data. This is useful for example when categorizing aggregate events.

For inference we use collapsed Gibbs Sampling (Griffiths and Steyvers, 2004) where each hidden variable, z_i , is sampled in turn, and parameters are integrated out. Example types are displayed in Figure 2. To estimate the distribution over types for a given event, a sample of the corresponding hidden variables is taken from the Gibbs markov chain after sufficient burn in. Prediction for new data is performed using a streaming approach to inference (Yao et al., 2009).

⁵This annotated data is provided in the paper’s supplementary materials.

Label	Top 5 Event Phrases	Top 5 Entities
Sports	tailgate - scrimmage - tailgating - homecoming - regular season	espn - ncaa - tigers - eagles - varsity
Concert	concert - presale - performs - concerts - tickets	taylor swift - toronto - britney spears - rihanna - rock
Perform	matinee - musical - priscilla - seeing - wicked	shrek - les mis - lee evans - wicked - broadway
TV	new season - season finale - finished season - episodes - new episode	jersey shore - true blood - glee - dvr - hbo
Movie	watch love - dialogue theme - inception - hall pass - movie	netflix - black swan - insidious - tron - scott pilgrim
Sports	inning - innings - pitched - homered - homer	mlb - red sox - yankees - twins - dl
Politics	presidential debate - osama - presidential candidate - republican debate - debate performance	obama - president obama - gop - cnn - america
TV	network news broadcast - airing - primetime drama - channel - stream	nbc - espn - abc - fox - mtv
Product	unveils - unveiled - announces - launches - wraps off	apple - google - microsoft - uk - sony
Meeting	shows trading - hall - mtg - zoning - briefing	town hall - city hall - club - commerce - white house
Finance	stocks - tumbled - trading report - opened higher - tumbles	reuters - new york - u.s. - china - euro
School	maths - english test - exam - revise - physics	english - maths - german - bio - twitter
Album	in stores - album out - debut album - drops on - hits stores	itunes - ep - uk - amazon - cd
TV	voted off - idol - scotty - idol season - dividend-paying	lady gaga - american idol - america - beyonce - glee
Religion	sermon - preaching - preached - worship - preach	church - jesus - pastor - faith - god
Conflict	declared war - war - shelling - opened fire - wounded	libya - afghanistan - #syria - syria - nato
Politics	senate - legislation - repeal - budget - election	senate - house - congress - obama - gop
Prize	winners - lotto results - enter - winner - contest	ipad - award - facebook - good luck - winners
Legal	bail plea - murder trial - sentenced - plea - convicted	casey anthony - court - india - new delhi - supreme court
Movie	film festival - screening - starring - film - gosling	hollywood - nyc - la - los angeles - new york
Death	live forever - passed away - sad news - condolences - buried	michael jackson - afghanistan - john lennon - young - peace
Sale	add into - 50% off - up - shipping - save up	groupon - early bird - facebook - @etsy - etsy
Drive	donate - tornado relief - disaster relief - donated - raise money	japan - red cross - joplin - june - africa

Figure 2: Example event types discovered by our model. For each type, the top 5 entities and event phrases are listed. The assigned labels are also presented.

Algorithm 1 Generative story for our data involving event types as hidden variables. Bayesian Inference techniques are applied to invert the generative process and infer an appropriate set of types to describe the observed events.

```

for each event type  $t = 1 \dots T$  do
  Generate  $\beta_t^n$  according to symmetric Dirichlet distribution  $\text{Dir}(\eta_n)$ .
  Generate  $\beta_t^d$  according to symmetric Dirichlet distribution  $\text{Dir}(\eta_d)$ .
end for
for each unique event phrase  $e = 1 \dots |E|$  do
  Generate  $\theta_e$  according to Dirichlet distribution  $\text{Dir}(\alpha)$ .
  for each entity which co-occurs with  $e$ ,  $i = 1 \dots N_e$  do
    Generate  $z_{e,i}^n$  from Multinomial( $\theta_e$ ).
    Generate the entity  $n_{e,i}$  from Multinomial( $\beta_{z_{e,i}^n}$ ).
  end for
  for each date which co-occurs with  $e$ ,  $i = 1 \dots N_d$  do
    Generate  $z_{e,i}^d$  from Multinomial( $\theta_e$ ).
    Generate the date  $d_{e,i}$  from Multinomial( $\beta_{z_{e,i}^d}$ ).
  end for
end for

```

4.1 Evaluation

To evaluate the ability of our model to classify events, we gathered 2 million extracted events of the form listed in Figure 1 (not including the type). We then ran Gibbs Sampling with 100 types for 1,000 iterations of burn-in, keeping the hidden variable assignments found in the last sample.

One of the authors manually inspected the resulting types and assigned them labels such as SPORTS, POLITICS, MUSICRELEASE and so on, based on their distribution over entities, and the event words which assign highest probability to that topic. In cases where the types appeared incoherent, they were labeled with a special symbol \emptyset . Note that this assignment of labels to types only needs to be done once and produces a labeling for an arbitrarily large number of event instances. Additionally the same set of types can easily be used to classify new event instances using streaming inference techniques (Yao et al., 2009). One interesting direction for future work is automatic labeling and coherence evaluation of automatically discovered event types analogous to recent work on topic models (Newman et al., 2010; Lau et al., 2011).

In order to evaluate the ability of our model to classify aggregate events, we grouped together all (entity,date) pairs in the data, then randomly sampled 100 of the 450 pairs with highest association

	Precision	Recall	F ₁
TwICAL-CLASSIFY	0.75	0.63	0.68
Supervised Baseline	0.69	0.41	0.52

Table 3: Precision and recall of event type categorization at the point of maximum F₁ score.

(see §5) for annotation using the event types discovered by our model. For prediction, we estimate the distribution over event types using streaming inference over the new data.

To help demonstrate the benefits of leveraging large quantities of unlabeled data for event classification, we compare against a supervised Maximum Entropy baseline which makes use of the 100 annotated events using 10-fold cross validation.⁶ For features, we treat the set of event phrases that co-occur with each (entity, date) pair as a bag-of-words. Because many event categories are infrequent, there are often few or no training examples for a category, leading to low performance.

Figure 3 compares the performance of our unsupervised approach to the supervised baseline, via a precision-recall curve obtained by varying the threshold on the probability of the most likely type. In addition table 3 compares precision and recall at the point of maximum F-score. Our unsupervised approach to event categorization achieves a 31% increase in F₁ over the supervised baseline.

5 Ranking Events

Simply using frequency to determine which events are significant is insufficient, because many tweets refer to common events in user’s daily lives. As an example, users often mention what they are eating for lunch, therefore entities such as *McDonalds* occur relatively frequently in association with references to most calendar days. Important events can be distinguished as those which have strong association with a unique date as opposed to being spread evenly across days on the calendar. To extract significant events of general interest from Twitter, we thus need some way to measure the strength of association between an entity and a date.

⁶As a fair comparison, the size of this training dataset is commensurate with the amount of effort required to label the automatically discovered types.




Figure 3: Precision and recall at predicting event types. Without any thresholding, the accuracy of the MaxEnt classifier is 0.42 and the accuracy of the Topic Modeling approach is 0.66 (these numbers represent starting points on the P/R curves). One of the highest-confidence examples predicted by our classifier happens to be incorrect in this dataset, this explains why the curve doesn’t reach 100% precision.

In order to measure the association strength between an entity and a specific date, we utilize the G^2 log likelihood ratio statistic. G^2 has been argued to be more appropriate for text analysis tasks than χ^2 (Dunning, 1993). Although Fisher’s Exact test would produce more accurate p-values (Moore, 2004), given the amount of data with which we are working (sample size greater than 10^{11}), it proves difficult to compute Fisher’s Exact Test Statistic, which results in floating point overflow even when using 64-bit operations. The G^2 test works sufficiently well in our setting, however, as computing association between entities and dates produces less sparse contingency tables than when working with pairs of entities (or words).

The G^2 test is based on the likelihood ratio between a model in which the entity is conditioned on the date, and a model of independence between entities and date references. For a given entity e and date d this statistic can be computed as follows:

$$G^2 = \sum_{x \in \{e, \neg e\}, y \in \{d, \neg d\}} O_{x,y} \times \ln \left(\frac{O_{x,y}}{E_{x,y}} \right)$$

Where $O_{e,d}$ is the observed fraction of tweets containing both e and d , $O_{e,-d}$ is the observed fraction of tweets containing e , but not d , and so on. Similarly $E_{e,d}$ is the expected fraction of tweets containing both e and d assuming a model of independence.

6 Experiments

To estimate the quality of the calendar entries generated using our approach we manually evaluated a sample of the top 100, 500 and 1,000 calendar entries occurring within a 2-week future window of November 3rd.

6.1 Data

For evaluation purposes, we gathered roughly the 100 million most recent tweets on November 3rd 2011 (collected using the Twitter Streaming API⁷, and tracking a broad set of temporal keywords, including “today”, “tomorrow”, names of weekdays, months, etc.).

We extracted named entities in addition to event phrases, and temporal expressions from the text of each of the 100M tweets. We then added the extracted triples to the dataset used for inferring event types described in §4, and performed 50 iterations of Gibbs sampling for predicting event types on the new data, holding the hidden variables in the original data constant. This *streaming* approach to inference is similar to that presented by Yao et al. (2009).

We then ranked the extracted events as described in §5, and randomly sampled 50 events from the top ranked 100, 500, and 1,000. We annotated the events with 4 separate criteria:

1. Is there a significant event involving the extracted entity which will take place on the extracted date?
2. Is the most frequently extracted event phrase informative?
3. Is the event’s type correctly classified?
4. Are each of (1-3) correct? That is, does the event contain a correct entity, date, event phrase, and type?

⁷<https://dev.twitter.com/docs/streaming-api>

Note that if (1) is marked as incorrect for a specific event, subsequent criteria are always marked incorrect.

6.2 Baseline

To demonstrate the importance of natural language processing and information extraction techniques in extracting informative events, we compare against a simple baseline which does not make use of the Ritter et. al. named entity recognizer or our event recognizer; instead, it considers all 1-4 grams in each tweet as candidate calendar entries, relying on the G^2 test to filter out phrases which have low association with each date.

6.3 Results

The results of the evaluation are displayed in table 4. The table shows the precision of the systems at different yield levels (number of aggregate events). These are obtained by varying the thresholds in the G^2 statistic. Note that the baseline is only comparable to the third column, i.e., the precision of (entity, date) pairs, since the baseline is not performing event identification and classification. Although in some cases ngrams do correspond to informative calendar entries, the precision of the ngram baseline is extremely low compared with our system.

In many cases the ngrams don’t correspond to salient entities related to events; they often consist of single words which are difficult to interpret, for example “Breaking” which is part of the movie “Twilight: Breaking Dawn” released on November 18. Although the word “Breaking” has a strong association with November 18, by itself it is not very informative to present to a user.⁸

Our high-confidence calendar entries are surprisingly high quality. If we limit the data to the 100 highest ranked calendar entries over a two-week date range in the future, the precision of extracted (entity, date) pairs is quite good (90%) - an 80% increase over the ngram baseline. As expected precision drops as more calendar entries are displayed, but re-

⁸In addition, we notice that the ngram baseline tends to produce many near-duplicate calendar entries, for example: “Twilight Breaking”, “Breaking Dawn”, and “Twilight Breaking Dawn”. While each of these entries was annotated as correct, it would be problematic to show this many entries describing the same event to a user.

November 2011						
Mon Nov 7	Tue Nov 8	Wed Nov 9	Thu Nov 10	Fri Nov 11	Sat Nov 12	Sun Nov 13
Justin <i>meet</i> Other	Paris <i>love</i> Other	EAS <i>test</i> Other	Robert Pattinson <i>show</i> Performance	iPhone <i>debut</i> Product Release	Sydney <i>perform</i> Other	Playstation <i>answers</i> Product Release
Motorola Pro+ <i>kick</i> Product Release	iPhone <i>holding</i> Product Release	The Feds <i>cut off</i> Other	James Murdoch <i>give evidence</i> Other	Remembrance Day <i>open</i> Performance	Pullman Ballroom <i>promoted</i> Other	Samsung Galaxy Tab <i>launch</i> Product Release
Nook Color 2 <i>launch</i> Product Release	Election Day <i>vote</i> Political Event	Toca Rivera <i>promoted</i> Performance	RTL-TVI <i>post</i> TV Event	France <i>play</i> Other	Fox <i>fight</i> Other	Sony <i>answers</i> Product Release
Eid-ul-Azha <i>celebrated</i> Performance	Blue Slide Park <i>listening</i> Music Release	Alert System <i>test</i> Other	Gotti Live <i>work</i> Other	Veterans Day <i>closed</i> Other	Plaza <i>party</i> Party	Chibi Chibi Burger other
MW3 <i>midnight release</i> Other	Hedley <i>album</i> Music Release	Max Day <i>give</i> Other	Bambi Awards <i>perform</i> Performance	Skyrim <i>arrives</i> Product Release	Red Carpet <i>invited</i> Party	Jiexpo Kemayoran <i>promoted</i> TV Event

Figure 4: Example future calendar entries extracted by our system for the week of November 7th. Data was collected up to November 5th. For each day, we list the top 5 events including the entity, *event phrase*, and *event type*. While there are several errors, the majority of calendar entries are informative, for example: the Muslim holiday eid-ul-azha, the release of several videogames: Modern Warfare 3 (MW3) and Skyrim, in addition to the release of the new playstation 3D display on Nov 13th, and the new iPhone 4S in Hong Kong on Nov 11th.

# calendar entries	precision				
	ngram baseline	entity + date	event phrase	event type	entity + date + event + type
100	0.50	0.90	0.86	0.72	0.70
500	0.46	0.66	0.56	0.54	0.42
1,000	0.44	0.52	0.42	0.40	0.32

Table 4: Evaluation of precision at different recall levels (generated by varying the threshold of the G^2 statistic). We evaluate the top 100, 500 and 1,000 (entity, date) pairs. In addition we evaluate the precision of the most frequently extracted event phrase, and the predicted event type in association with these calendar entries. Also listed is the fraction of cases where all predictions (“entity + date + event + type”) are correct. We also compare against the precision of a simple ngram baseline which does not make use of our NLP tools. Note that the ngram baseline is only comparable to the entity+date precision (column 3) since it does not include event phrases or types.

mains high enough to display to users (in a ranked list). In addition to being less likely to come from extraction errors, highly ranked entity/date pairs are more likely to relate to popular or important events, and are therefore of greater interest to users.

In addition we present a sample of extracted future events on a calendar in figure 4 in order to give an example of how they might be presented to a user. We present the top 5 entities associated with each date, in addition to the most frequently extracted event phrase, and highest probability event type.

6.4 Error Analysis

We found 2 main causes for why entity/date pairs were uninformative for display on a calendar, which occur in roughly equal proportion:

Segmentation Errors Some extracted “entities” or

ngrams don’t correspond to named entities or are generally uninformative because they are mis-segmented. Examples include “RSVP”, “Breaking” and “Yikes”.

Weak Association between Entity and Date

In some cases, entities are properly segmented, but are uninformative because they are not strongly associated with a specific event on the associated date, or are involved in many different events which happen to occur on that day. Examples include locations such as “New York”, and frequently mentioned entities, such as “Twitter”.

7 Related Work

While we are the first to study open domain event extraction within Twitter, there are two key related

strands of research: extracting specific types of events from Twitter, and extracting open-domain events from news (Pustejovsky et al., 2003).

Recently there has been much interest in information extraction and event identification within Twitter. Benson et al. (2011) use distant supervision to train a relation extractor which identifies artists and venues mentioned within tweets of users who list their location as New York City. Sakaki et al. (2010) train a classifier to recognize tweets reporting earthquakes in Japan; they demonstrate their system is capable of recognizing almost all earthquakes reported by the Japan Meteorological Agency. In contrast to previous work on Twitter event identification, our approach is independent of event type or domain and is thus more widely applicable. Additionally, our work focuses on extracting a calendar of events, extracting event-referring expressions and categorizing events into types.

Also relevant is work on identifying events (Kumaran and Allan, 2004; Das Sarma et al., 2011; Bethard and Martin, 2006), and extracting timelines (Ling and Weld, 2010) from news articles.⁹ Twitter status messages present both unique challenges and opportunities when compared with news articles. Twitter’s noisy text presents serious challenges for NLP tools. On the other hand, it contains a higher proportion of references to present and future dates. Tweets do not require complex reasoning about relations between events in order to place them on a timeline as is typically necessary in long texts containing narratives (Song and Cohen, 1991). Additionally, unlike News, Tweets often discuss mundane events which are not of general interest, so it is crucial to exploit redundancy of information to assess whether an event is significant.

Previous work on open-domain information extraction (Banko et al., 2007; Weld et al., 2009; Fader et al., 2011) has mostly focused on extracting relations (as opposed to events) from web corpora and has also extracted relations based on verbs. In contrast, this work extracts on events, using tools adapted to Twitter’s noisy text, and extracts event phrases which are often adjectives or nouns, for example: *Super Bowl Party on Feb 5th*.

8 Conclusions

We have presented a scalable and open-domain approach to extracting and categorizing events from status messages. We evaluated the quality of these events in a manual evaluation showing a clear improvement in performance over an ngram baseline

We proposed a novel approach to categorizing events in an open-domain text genre with unknown types. Our approach based on latent variable models first discovers event types which match the data, which are then used to classify aggregate events without any annotated examples. Because this approach is able to leverage large quantities of unlabeled data, it outperforms a supervised baseline by 31%.

A possible avenue for future work is extraction of even richer event representations, while maintaining domain independence. For example: grouping together related entities, classifying entities in relation to their roles in the event, thereby, extracting a frame-based representation of events.

References

- James Allan, Ron Papka, and Victor Lavrenko. 1998. On-line new event detection and tracking. In *SIGIR*.
- Michele Banko, Michael J Cafarella, Stephen Soderl, Matt Broadhead, and Oren Etzioni. 2007. Open information extraction from the web. In *IJCAI*.
- Edward Benson, Aria Haghighi, and Regina Barzilay. 2011. Event discovery in social media feeds. In *ACL*.
- Steven Bethard and James H. Martin. 2006. Identification of event mentions and their semantic class. In *EMNLP*.
- Nathanael Chambers and Dan Jurafsky. 2011. Template-based information extraction without the templates. In *Proceedings of ACL*.
- Nathanael Chambers, Shan Wang, and Dan Jurafsky. 2007. Classifying temporal relations between events. In *ACL*.
- Anish Das Sarma, Alpa Jain, and Cong Yu. 2011. Dynamic relationship and event discovery. In *WSDM*.
- G. Doddington, A. Mitchell, M. Przybocki, L. Ramshaw, S. Strassel, and R. Weischedel. 2004. The Automatic Content Extraction (ACE) Program—Tasks, Data, and Evaluation. *LREC*.
- Ted Dunning. 1993. Accurate methods for the statistics of surprise and coincidence. *Comput. Linguist*.

⁹<http://newstimeline.googlelabs.com/>

- Elena Erosheva, Stephen Fienberg, and John Lafferty. 2004. Mixed-membership models of scientific publications. *PNAS*.
- Oren Etzioni. 1996. Moving up the information food chain: deploying softbots on the world wide web. In *Proceedings of the thirteenth national conference on Artificial intelligence - Volume 2*.
- Anthony Fader, Stephen Soderland, and Oren Etzioni. 2011. Identifying relations for open information extraction. In *EMNLP*.
- Kevin Gimpel, Nathan Schneider, Brendan O'Connor, Dipanjan Das, Daniel Mills, Jacob Eisenstein, Michael Heilman, Dani Yogatama, Jeffrey Flanigan, and Noah A. Smith. 2011. Part-of-speech tagging for twitter: Annotation, features, and experiments. In *ACL*.
- T. L. Griffiths and M. Steyvers. 2004. Finding scientific topics. *Proc Natl Acad Sci U S A*, 101 Suppl 1.
- Ralph Grishman and Beth Sundheim. 1996. Message understanding conference - 6: A brief history. In *Proceedings of the International Conference on Computational Linguistics*.
- Zornitsa Kozareva and Eduard Hovy. 2010. Learning arguments and supertypes of semantic relations using recursive patterns. In *ACL*.
- Giridhar Kumaran and James Allan. 2004. Text classification and named entities for new event detection. In *SIGIR*.
- John D. Lafferty, Andrew McCallum, and Fernando C. N. Pereira. 2001. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *ICML*.
- Jey Han Lau, Karl Grieser, David Newman, and Timothy Baldwin. 2011. Automatic labelling of topic models. In *ACL*.
- Xiao Ling and Daniel S. Weld. 2010. Temporal information extraction. In *AAAI*.
- Xiaohua Liu, Shaodian Zhang, Furu Wei, and Ming Zhou. 2011. Recognizing named entities in tweets. In *ACL*.
- Inderjeet Mani and George Wilson. 2000. Robust temporal processing of news. In *ACL*.
- Inderjeet Mani, Marc Verhagen, Ben Wellner, Chong Min Lee, and James Pustejovsky. 2006. Machine learning of temporal relations. In *ACL*.
- Robert C. Moore. 2004. On log-likelihood-ratios and the significance of rare events. In *EMNLP*.
- David Newman, Jey Han Lau, Karl Grieser, and Timothy Baldwin. 2010. Automatic evaluation of topic coherence. In *HLT-NAACL*.
- Diarmuid Ó Séaghdha. 2010. Latent variable models of selectional preference. In *ACL, ACL '10*.
- J. Pustejovsky, P. Hanks, R. Sauri, A. See, R. Gaizauskas, A. Setzer, D. Radev, B. Sundheim, D. Day, L. Ferro, and M. Lazo. 2003. The TIMEBANK corpus. In *Proceedings of Corpus Linguistics 2003*.
- Alan Ritter, Mausam, and Oren Etzioni. 2010. A latent dirichlet allocation method for selectional preferences. In *ACL*.
- Alan Ritter, Sam Clark, Mausam, and Oren Etzioni. 2011. Named entity recognition in tweets: An experimental study. *EMNLP*.
- Takeshi Sakaki, Makoto Okazaki, and Yutaka Matsuo. 2010. Earthquake shakes twitter users: real-time event detection by social sensors. In *WWW*.
- Roser Sauri, Robert Knippen, Marc Verhagen, and James Pustejovsky. 2005. Evita: a robust event recognizer for qa systems. In *HLT-EMNLP*.
- Fei Song and Robin Cohen. 1991. Tense interpretation in the context of narrative. In *Proceedings of the ninth National conference on Artificial intelligence - Volume 1, AAAI'91*.
- Benjamin Van Durme and Daniel Gildea. 2009. Topic models for corpus-centric knowledge generalization. In *Technical Report TR-946, Department of Computer Science, University of Rochester, Rochester*.
- Daniel S. Weld, Raphael Hoffmann, and Fei Wu. 2009. Using wikipedia to bootstrap open information extraction. *SIGMOD Rec*.
- Yiming Yang, Tom Pierce, and Jaime Carbonell. 1998. A study of retrospective and on-line event detection. In *Proceedings of the 21st annual international ACM SIGIR conference on Research and development in information retrieval, SIGIR '98*.
- Limin Yao, David Mimno, and Andrew McCallum. 2009. Efficient methods for topic model inference on streaming document collections. In *KDD*.
- Limin Yao, Aria Haghighi, Sebastian Riedel, and Andrew McCallum. 2011. Structured relation discovery using generative models. In *EMNLP*.
- Fabio Massimo Zanzotto, Marco Pennacchiotti, and Kostas Tsioutsoulis. 2011. Linguistic redundancy in twitter. In *EMNLP*.